


A Neural Networks Presentation

Lawrence Chan

An out-of-date joke




IN CS, IT CAN BE HARD TO EXPLAIN
THE DIFFERENCE BETWEEN THE EASY
AND THE VIRTUALLY IMPOSSIBLE.

What is Machine Learning?

Traditional Programming




Machine Learning




Deep Neural Networks

- Complicated function made of simple parts
 - “Universal approximator property”
- A lot of recent success in ML/AI has been due to better architectures for and increased ability to train neural nets
- Very messy; not a lot of good theory behind them.



A “Neuron”




Train a Network by Minimizing a Loss Function

- Loss functions measure how well the classifier is doing
- For example:
 - 0-1 Loss
 - Word-Edit Distance
 - Mean-squared error
 - Cross-Entropy Loss
- Warning: loss functions not guaranteed to capture what you want!

Some math on the board.

Modern Deep Learning - “Programming by gradient descent”



Requires **differentiable** loss function!

Picture Credit: Andrew Ng

Backpropagation - Efficiently computing gradients

- Chain rule!

$$\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx},$$


More math on board

Thankfully, you don't actually need any of that to do (simple) deep learning!

- Machine Learning APIs
 - Google Cloud Machine Learning
 - Azure Machine Learning
- High level interfaces (handle almost everything):
 - Sklearn
 - Keras
- Lower-level interfaces (still handle backprop and gradient updates)
 - Tensorflow
 - PyTorch
- We're just going to use Scikit-Learn.

One last warning: overfitting!

- Model may be fitting to noise instead of capturing “true” trend.
- Ways to address:
 - Use unseen data to measure how good your classifier is “really” performing.
 - Regularization
 - Visualizing Decision Boundary (aka, manual inspection)



A Pause for Questions

Onto a practical demonstration!

Available at github.com/chanlaw/nn-presentation